

## Sarsa: On-Policy TD Control

```


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
    Initialize  $S$ 
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
         $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$ 
         $S \leftarrow S'; A \leftarrow A'$ ;
    until  $S$  is terminal

```

## Q-Learning: Off-Policy TD Control

```


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
    Initialize  $S$ 
    Repeat (for each step of episode):
        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
        Take action  $A$ , observe  $R, S'$ 
         $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
         $S \leftarrow S'$ ;
    until  $S$  is terminal

```

**Algorithm 1** Simulated annealing optimisation method. **Algorithm 2** Tabu search optimisation method.

**Require:** Input( $T_0, \alpha, N, T_f$ )

- 1:  $T \leftarrow T_0$
- 2:  $S_{act} \leftarrow$  generate initial solution
- 3: **while**  $T \geq T_f$  **do**
- 4:   **for** cont  $\leftarrow 1$  TO  $N(T)$  **do**
- 5:      $S_{cond} \leftarrow$  Neighbour solution [from ( $S_{act}$ )]
- 6:      $\delta \leftarrow f(S_{cond}) - f(S_{act})$
- 7:     **if** rand(0,1)  $< e^{-\delta/T}$  **or**  $\delta < 0$  **then**
- 8:        $S_{act} \leftarrow S_{cond}$
- 9:     **end if**
- 10:   **end for**
- 11:    $T \leftarrow \alpha(T)$
- 12: **end while**
- 13: **return** Best  $S_{act}$  visited

**Algorithm 3** Genetic algorithm optimisation method.


- 1:  $t \leftarrow 0$
- 2: Initialise  $P(t)$  ▶ initial population
- 3: Evaluate  $P(t)$
- 4: **repeat**
- 5:   Generate offspring  $C(t)$  from  $P(t)$  ▶ using crossover and mutation
- 6:   Evaluate  $C(t)$
- 7:   Select  $P(t+1)$  from  $P(t) \cup C(t)$
- 8:    $t \leftarrow t + 1$
- 9: **until** a termination criterion is satisfied

**Cumulative Distribution Function (CDF) of pixels in image X and Y:**

$$cdf_x(i) = \sum_{j=0}^i p_x(j)$$

$$cdf_y(i) = (i+1)K, \text{ for } 0 \leq i < L \text{ for some constant } K.$$

$$P(\sim A|B) = 1 - P(A|B)$$



$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D)$$

$$\text{TF}(t, d) = \frac{\text{Number of times term "t" appears in document d}}{\text{Total number of terms in document d}}$$

$$\text{IDF}(t, D) = \log_{10} \frac{\text{Total number of documents in corpus D}}{\text{Number of documents containing term t}}$$

**Algorithm 1** Memory-based explainable reinforcement learning approach with the on-policy method SARSA to compute the probability of success and the number of transitions to the goal state.

- 1: Initialize  $Q(s, a), T_s, T_s, P_s, N_t$
- 2: **for** each episode **do**
- 3:   Initialize  $T_{List}[]$
- 4:   Choose an action using  $a_t \leftarrow \text{SELECTACTION}(s_t)$
- 5:   **repeat**
- 6:     Take action  $a_t$
- 7:     Save state-action transition  $T_{List}.\text{add}(s, a)$
- 8:      $T_s[s][a] \leftarrow T_s[s][a] + 1$
- 9:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$
- 10:    Choose next action  $a_{t+1}$  using softmax action selection method
- 11:     $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$
- 12:     $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$
- 13: **until**  $s_t$  is terminal (goal or aversive state)
- 14: **if**  $s_t$  is goal state **then**
- 15:    **for** each  $s, a \in T_{List}$  **do**
- 16:        $T_s[s][a] \leftarrow T_s[s][a] + 1$
- 17:    **end for**
- 18: **end if**
- 19: Compute  $P_s \leftarrow T_s/T_t$
- 20: Compute  $N_t$  for each  $s \in T_{List}$  as  $\text{pos}(s, T_{List}) + 1$
- 21: **end for**

**Algorithm 2** Explainable reinforcement learning approach to compute the probability of success using the learning-based approach.

- 1: Initialize  $Q(s, a), \mathbb{P}(s_t, a_t)$
- 2: **for** each episode **do**
- 3:   Initialize  $s_t$
- 4:   Choose an action  $a_t$  from  $s_t$
- 5:   **repeat**
- 6:     Take action  $a_t$
- 7:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$
- 8:     Choose next action  $a_{t+1}$  using softmax action selection method
- 9:      $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$
- 10:     $\mathbb{P}(s_t, a_t) \leftarrow \mathbb{P}(s_t, a_t) + \alpha[\varphi_{t+1} + \mathbb{P}(s_{t+1}, a_{t+1}) - \mathbb{P}(s_t, a_t)]$
- 11:     $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$
- 12: **until**  $s_t$  is terminal (goal or aversive state)

**Algorithm 3** Explainable reinforcement learning approach to compute the probability of success using introspection-based approach.

- 1: Initialize  $Q(s, a), \hat{P}_s$
- 2: **for** each episode **do**
- 3:   Initialize  $s_t$
- 4:   Choose an action  $a_t$  from  $s_t$
- 5:   **repeat**
- 6:     Take action  $a_t$
- 7:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$
- 8:     Choose next action  $a_{t+1}$  using softmax action selection method
- 9:      $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$
- 10:     $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$
- 11: **until**  $s_t$  is terminal (goal or aversive state)
- 12:  $\hat{P}_s \approx \left[ (1 - \sigma) \cdot \left( \frac{1}{2} \cdot \log_{10} \frac{Q(s_t, a_t)}{R^t} + 1 \right) \right]_{\hat{P}_s \leq 1}$

## Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS) calls DFS for different depths

Uniform Cost Search (UCS)

expand least-cost unexpanded node (按照所有的距离累加排序直到目标)

Greedy Best-First Search priority queue (base heuristic) 按照h从小到大

A\* Search  $g(n) + h(n)$  (从小到大)

## Ensemble Learning, Random Forests:

```

1 Select the number of models to build, m
2 for i = 1 to m do
3   Generate a bootstrap sample of the original data
4   Train a tree model on this sample
5   for each split do
6     Randomly select k (< P) of the original predictors
7     Select the best predictor among the k predictors and
8     partition the data
9   end
10  Use typical tree model stopping criteria to determine when a
    tree is complete (but do not prune)
11 end

```


## Boosting:

1. Initialize the observation weights  $w_i = 1/N$ ,  $i=1,2,\dots,N$ .
2. For  $m=1$  to  $M$ :
  - (a) Fit a classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - (b) Compute
 
$$\text{err}_m = \frac{\sum_1^N w_i I(y_i \neq G_m(x_i))}{\sum_1^N w_i}$$
  - (c) Compute the influence  $\alpha_m$ 

正确: 乘  $e^{-\alpha}$   
错误: 乘  $e^{+\alpha}$

$$\alpha_m = \frac{1}{2} \ln \frac{1 - \text{err}_m}{\text{err}_m}$$
  - (d) Set  $w_i^{\text{new\_iteration}} = w_i^{\text{old\_iteration}} e^{\pm \alpha}$ ,  $i=1,2,\dots,N$ .
3. Output  $G(x) = \text{sign}[\sum_{m=1}^M \alpha_m G_m(x)]$ .

3. **Backward pass:** Propagate errors back through the network to adjust weights ( $E = \frac{1}{2} \sum_i (z_i - t_i)^2$ ).



$$X_n = (X - X_{\min}) / (X_{\max} - X_{\min}); X_n \in [0,1] \text{ or}$$

$$X_n = 2 * (X - X_{\min}) / (X_{\max} - X_{\min}) - 1; X_n \in [-1,1]$$

Rule of thumb:  $N_h$  should lead to a number of parameters (weights)  $N_w$  that:

$$N_w < (\text{Number of samples}) / 10$$

The number of weights  $N_w$  of a MLP, with  $N_i$  neurons in its input layer, a hidden layer with  $N_h$  neurons, and  $N_o$  neurons in the output layer is:

$$N_w = (N_i + 1) * N_h + (N_h + 1) * N_o$$

Machine learning (subfields of AI algorithms)

three types: supervised learning

unsupervised learning

reinforcement learning

① supervised learning: { Regressor (one-output, real value)

Binary classification (two discrete classes (positive/negative))

Multiclass classification (discrete classes, > 2 possible values)

real world inputs

Model inputs

Model

Model outputs

real world outputs

method: decision tree Entropy  $H(p_1, \dots, p_n) = \sum_{i=1}^n p_i \log_2 p_i$

Minimal error Pruning:  $E = 1 - \frac{n_{\text{node}}}{N + k} \rightarrow \text{number of class items}$

② unsupervised learning

learn about a dataset without labels

clustering: Grouping similar data point together

neural network

mathematical formula:  $Z = g(s) = w_0 + \sum_i w_i x_i$  (transfer function)


perception: neuron with step transfer function (like 0, 1)

and ( $w_0 = -1.5$ ,  $w_1 = 1$ ,  $w_2 = 1$ )

or ( $w_0 = -0.5$ ,  $w_1 = 1$ ,  $w_2 = 1$ )

learning:

error function:  $E = \frac{1}{2} \sum_i (Z_i - t_i)^2$



$$\text{if } Z(s) = \frac{1}{1 + e^{-s}}, Z'(s) = Z(1 - Z)$$

$$Z(s) = \tanh(s), Z'(s) = 1 - Z^2$$

step: Forward pass : hit all the formulas

$$\text{Calculate the error: } E = \frac{1}{2} \sum_i (Z_i - t_i)^2$$

Backward pass : adjust weight we derivative

## Returns

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- $V^\pi(s)$  is the expected value of following policy  $\pi$  in state  $s$
- $V^*(s)$  be the maximum discounted reward obtainable from  $s$

$$Q(s, a) = r(s, a) + \gamma V^*(s')$$

$$O(1) < O(\log n) < O(n) < O(n \log n) < O(n^2) < O(n^3) < O(2^n)$$