

## Sarsa: On-Policy TD Control

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
 Repeat (for each episode):  
 Initialize  $S$   
 Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
 Repeat (for each step of episode):  
 Take action  $A$ , observe  $R, S'$   
 Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$$S \leftarrow S'; A \leftarrow A';$$

$$\epsilon \leftarrow \epsilon - \epsilon_{\text{decay}}$$

**Algorithm 1** Simulated annealing optimisation method.

**Require:** Input( $T_0, \alpha, N, T_f$ )  
 1:  $T \leftarrow T_0$   
 2:  $S_{act} \leftarrow$  generate initial solution  
 3: **while**  $T \geq T_f$  **do**  
 4:   **for** cont  $\leftarrow 1$  TO  $N(T)$  **do**  
 5:      $S_{cond} \leftarrow$  Neighbour solution [from ( $S_{act}$ )]  
 6:      $\delta \leftarrow f(S_{cond}) - f(S_{act})$   
 7:     **if** rand(0,1)  $< e^{-\delta/T}$  **or**  $\delta < 0$  **then**  
 8:        $S_{act} \leftarrow S_{cond}$   
 9:     **end if**  
 10:   **end for**  
 11:    $T \leftarrow \alpha(T)$   
 12: **end while**  
 13: **return** Best  $S_{act}$  visited

**Algorithm 3** Genetic algorithm optimisation method.

1:  $t \leftarrow 0$   
 2: Initialise  $P(t)$  ▶ initial population  
 3: Evaluate  $P(t)$   
 4: **repeat**  
 5:   Generate offspring  $C(t)$  from  $P(t)$  ▶ using crossover and mutation  
 6:   Evaluate  $C(t)$   
 7:   Select  $P(t+1)$  from  $P(t) \cup C(t)$   
 8:    $t \leftarrow t + 1$   
 9: **until** a termination criterion is satisfied  
 10: **return** Best individual found from  $P$

**Cumulative Distribution Function (CDF) of pixels in image X and Y:**

$$cdf_X(i) = \sum_{j=0}^i p_x(j)$$

$$cdf_Y(i) = (i+1)K, \text{ for } 0 \leq i < L \text{ for some constant } K.$$

## Q-Learning: Off-Policy TD Control

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
 Repeat (for each episode):  
 Initialize  $S$   
 Repeat (for each step of episode):  
 Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
 Take action  $A$ , observe  $R, S'$   

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$$S \leftarrow S';$$

$$\epsilon \leftarrow \epsilon - \epsilon_{\text{decay}}$$

**Algorithm 2** Tabu search optimisation method.

1:  $s_0 \leftarrow$  generate initial solution  
 2:  $s_{best} \leftarrow s_0$   
 3:  $\text{tabuList} \leftarrow \{s_0\}$   
 4: **repeat**  
 5:    $\{s_1, s_2, \dots, s_n\} \leftarrow$  generate neighbourhood from ( $s_{best}$ )  
 6:    $s_{candidate} \leftarrow s_1$   
 7:   **for**  $i \leftarrow 2$  TO  $n$  **do**  
 8:      $\delta \leftarrow f(s_i) - f(s_{candidate})$   
 9:     **if**  $s_i$  is not in  $\text{tabuList}$  **and**  $\delta < 0$  **then**  
 10:        $s_{candidate} \leftarrow s_i$   
 11:     **end if**  
 12:   **end for**  
 13:    $s_{best} \leftarrow s_{candidate}$   
 14:   Add  $s_{candidate}$  to  $\text{tabuList}$   
 15: **until** a termination criterion is satisfied

• deduction: cause + rule  $\Rightarrow$  effect

• abduction: effect + rule  $\Rightarrow$  cause

• induction: cause + effect  $\Rightarrow$  rule

A production rule has the form

if <condition> then <conclusion>

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

$$P(A) = P(A \wedge B) + P(A \wedge \neg B)$$

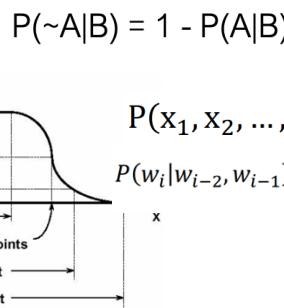
$$P(A|B) = \frac{P(A, B)}{P(B)}$$

$$P(A, B, C, D) = P(A|B, C, D)P(B|C, D)P(C|D)P(D)$$

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\neg A)P(\neg A)}$$

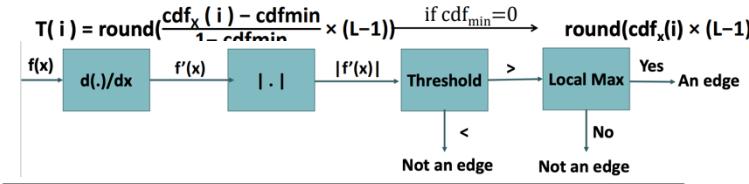
$$P(x_1, x_2, \dots, x_n|y) = P(x_1|y) \cdot P(x_2|y) \cdots P(x_n|y)$$

$$P(w_i|w_{i-2}, w_{i-1}) = \frac{P(w_i, w_{i-2}, w_{i-1})}{P(w_{i-2}, w_{i-1})} \approx \frac{\text{count}(w_i, w_{i-2}, w_{i-1})}{\text{count}(w_{i-2}, w_{i-1})}$$



$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D)$$

So, the transfer function can be defined as:



$$\text{TF}(t, d) = \frac{\text{Number of times term "t" appears in document d}}{\text{Total number of terms in document d}}$$

$$\text{IDF}(t, D) = \log_{10} \frac{\text{Total number of documents in corpus D}}{\text{Number of documents containing term t}}$$

**Algorithm 1** Memory-based explainable reinforcement learning approach with the on-policy method SARSA to compute the probability of success and number of transitions to the goal state.

1: Initialize  $Q(s, a), T_t, T_s, P_s, N_t$   
 2: **for** each episode **do**  
 3:   Initialize  $T_{List}[]$   
 4:   Choose an action using  $a_t \leftarrow \text{SELECTACTION}(s_t)$   
 5:   **repeat**  
 6:     Take action  $a_t$   
 7:     Save state-action transition  $T_{List}.\text{add}(s, a)$   
 8:      $T_t[s][a] \leftarrow T_t[s][a] + 1$   
 9:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$   
 10:   Choose next action  $a_{t+1}$  using softmax action selection method  
 11:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$   
 12:    $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$   
 13:   **until**  $s$  is terminal (goal or aversive state)  
 14:   **if**  $s$  is goal state **then**  
 15:     **for** each  $s, a \in T_{List}$  **do**  
 16:        $T_s[s][a] \leftarrow T_s[s][a] + 1$   
 17:     **end for**  
 18:   **end if**  
 19:   Compute  $P_s \leftarrow T_s/T_t$   
 20:   Compute  $N_t$  for each  $s \in T_{List}$  as  $\text{pos}(s, T_{List}) + 1$   
 21: **end for**

**Algorithm 2** Explainable reinforcement learning approach to compute the probability of success using the learning-based approach.

1: Initialize  $Q(s, a), \mathbb{P}(s_t, a_t)$   
 2: **for** each episode **do**  
 3:   Initialize  $s_t$   
 4:   Choose an action  $a_t$  from  $s_t$   
 5:   **repeat**  
 6:     Take action  $a_t$   
 7:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$   
 8:     Choose next action  $a_{t+1}$  using softmax action selection method  
 9:      $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$   
 10:      $\mathbb{P}(s_t, a_t) \leftarrow \mathbb{P}(s_t, a_t) + \alpha [\varphi_{t+1} + \mathbb{P}(s_{t+1}, a_{t+1}) - \mathbb{P}(s_t, a_t)]$   
 11:      $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$   
 12:   **until**  $s_t$  is terminal (goal or aversive state)

**Algorithm 3** Explainable reinforcement learning approach to compute the probability of success using the introspection-based approach.

1: Initialize  $Q(s, a), \hat{P}_s$   
 2: **for** each episode **do**  
 3:   Initialize  $s_t$   
 4:   Choose an action  $a_t$  from  $s_t$   
 5:   **repeat**  
 6:     Take action  $a_t$   
 7:     Observe reward  $r_{t+1}$  and next state  $s_{t+1}$   
 8:     Choose next action  $a_{t+1}$  using softmax action selection method  
 9:      $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$   
 10:      $\hat{P}_s \leftarrow \hat{P}_s + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$   
 11:      $s_t \leftarrow s_{t+1}; a_t \leftarrow a_{t+1}$   
 12:   **until**  $s_t$  is terminal (goal or aversive state)

On-Policy

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Off-Policy

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

## Search Algorithms Covered at this Lecture

### Breath First Search (BFS)

### Depth First Search (DFS)

### Iterative Deepening Search (IDDFS) calls DFS for different depths

### Uniform Cost Search (UCS)

expand least-cost unexpanded node (按照所有的距离累加排序直到目标)

### Greedy Best-First Search

priority queue (base heuristic)按照h从小到大

### A\* Search

g(n) + h(n) (从小到大)

Machine learning (subfields of AI algorithms)

three types: supervised learning

unsupervised learning

reinforcement learning

① supervised learning:

Regressor (one-output, real value)  
Binary classification (two discrete classes (positive/negative))  
Multi-class classification (discrete classes, > 2 possible values)

real world inputs Model inputs Model Model output real world outputs

method: decision tree Entropy  $H(p_1, \dots, p_n) = \sum_{i=1}^n p_i \log_2 p_i$

Minimal error Pruning:  $E = 1 - \frac{n_{\text{true}}}{N+K} \rightarrow \frac{\text{number in majority class}}{\text{total items}}$

② unsupervised learning

learn about a dataset without labels

clustering: Grouping similar data points together

neural network

mathematical formula:  $Z = g(s) = w_0 + \sum_i w_i x_i$  weight ↑  
↓ bias (constant) input (transfer function)

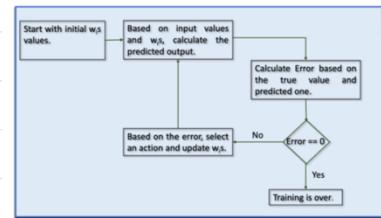
perception: neuron with step transfer function (like 0, 1)

and ( $w_0 = -1.5, w_1 = 1, w_2 = 1$ )

or ( $w_0 = -0.5, w_1 = 1, w_2 = 1$ )

learning:

error function:  $E = \frac{1}{2} \sum_i (Z_i - t_i)^2$  actual output ↓ target output



$$\text{if } Z(s) = \frac{1}{1+e^{-s}}, Z'(s) = Z(1-Z)$$

$$Z(s) = \tanh(s), Z'(s) = 1 - Z^2$$

step: Forward pass: list all the formulas

$$\text{Calculate the error: } E = \frac{1}{2} \sum_i (Z_i - t_i)^2$$

Backward pass: adjust weight via derivative

## Returns

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

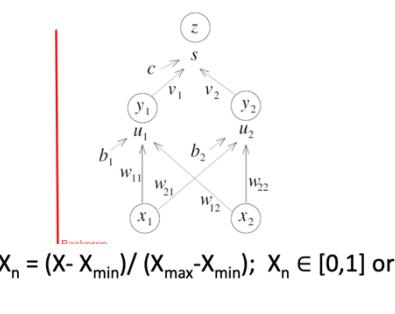
- $V^\pi(s)$  is the expected value of following policy  $\pi$  in state  $s$
- $V^*(s)$  be the maximum discounted reward obtainable from  $s$

$$Q(s, a) = r(s, a) + \gamma V^*(s')$$

$$O(1) < O(\log n) < O(n) < O(n \log n) < O(n^2) < O(n^3) < O(2^n)$$

3. **Backward pass:** Propagate errors back through the network to adjust weights ( $E = \frac{1}{2} \sum_i (Z_i - t_i)^2$ ).

$$\begin{aligned} \text{Partial derivative:} \\ \frac{\partial E}{\partial z} &= z - t \\ \frac{\partial z}{\partial s} &= g'(s) = z(1-z) \\ \frac{\partial s}{\partial y_1} &= v_1 \\ \frac{\partial y_1}{\partial u_1} &= y_1(1-y_1) \\ \text{Then: } w_{ij}^{\text{new\_epoch}} &\leftarrow w_{ij}^{\text{old\_epoch}} - \eta \frac{\partial E}{\partial w_{ij}} \\ \frac{\partial E}{\partial s} &= (z-t)z(1-z) \\ \frac{\partial s}{\partial E} &= (z-t)z(1-z)y_1 \\ \frac{\partial y_1}{\partial v_1} &= (z-t)z(1-z)y_1 \\ \frac{\partial v_1}{\partial u_1} &= \frac{\partial E}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial y_1} \frac{\partial y_1}{\partial u_1} = (z-t)z(1-z)v_1y_1(1-y_1) \\ \frac{\partial E}{\partial w_{11}} &= (z-t)z(1-z)v_1y_1(1-y_1)x_1 \end{aligned}$$



Rule of thumb:  $N_h$  should lead to a number of parameters (weights)  $N_w$  that:

$$N_w < (\text{Number of samples}) / 10$$

The number of weights  $N_w$  of a MLP, with  $N_i$  neurons in its input layer, a hidden layer with  $N_h$  neurons, and  $N_o$  neurons in the output layer is:

$$N_w = (N_i + 1) * N_h + (N_h + 1) * N_o$$