

Search

COMP3411: Artificial Intelligence



UNSW
SYDNEY


Search

Path 1:
Cost \$15, takes 20 mins


Path 2:
Cost \$20, takes 15 mins

Path 3:
Cost \$20, takes 20 mins, but you can pick
up your friend living in 21st St.

Which one is the best path?



Search



What is Search and why it is needed

“

In which we see how an agent can find a sequence of actions that achieves its goals when no single action will do.

- Peter Norvig and Stuart J. Russell
Artificial Intelligence: A Modern Approach

”

When is Search Needed?

Motion Planning



Navigation

Speech and Natural Language

Task Planning

Machine Learning

Game Playing



Converting real-world problem into machine
understandable problem

Let's formalize it!

A problem can be described by

- state (S)

2	3	5
8	6	1
7	4	

$[[2,3,5],$
 $[8,6,1],$
 $[7,4,0]]$

- action (A) $\{\text{"up"}, \text{"left"}\}$

- Transition Function do: $S \times A \rightarrow S$

Let's formalize it


Transition Function





Let's formalize it

Special states:

- Initial state: s_0




- Terminal states: $Z \subset S$




State Space in Graph

- Node in Graph is a data structure which contains a state and other related information such as legal actions from the state



State Space in Tree

- Node in Tree is a data structure which contains a state and other related information such as **children nodes**, **parent node** and **depth**



Uninformed Searches

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)

Uniform Cost Search (UCS)


Greedy Best-First Search

A* Search

Problem Solving by Graph Searching

- **Expanded nodes:** green and black
- **Frontiers:** red
- **Generated:** green, black and red
- **Not yet generated nodes:** blue

Search strategy differ in the way they expand the frontier



Uninformed Search

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)


Uniform Cost Search (UCS)

Greedy Best-First Search


A* Search

Informed Searches

Breath First Search (BFS)



Breadth-first Search Frontier



Breadth-First Search

- ✓ All nodes are expanded at a same depth in the tree before any nodes at the next level are expanded
- ✓ Can be implemented by using a queue to store frontier node
 - ✓ put newly generated successors at end of queue
- ✓ Finds the shallowest path first

Uninformed Search

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)


Uniform Cost Search (UCS)

Greedy Best-First Search


A* Search

Informed Searches

Depth First Search (DFS)



Depth-first Search - DFS




Depth First Search

Expand one node at the deepest level reached so far

Implementation:


- Implement the frontier as a stack, i.e. insert newly generated states at the front of the open list (frontier)
- Can be implemented by recursive function calls ✓

Depth-first Search Example



Depth First Search

- At any point depth-first search stores single path from root to leaf.
- Stop when node with goal state is expanded
- Include check that state has not already been explored **along a path – cycle checking**



Uninformed Searches

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)


Uniform Cost Search (UCS)

Greedy Best-First Search

A* Search

Iterative Deepening Depth-First Search (IDDFS)

- IDDFS calls DFS for different depths starting from an initial value. In every call, DFS is restricted from going beyond given depth.




- Iteration 0: A
- Iteration 1: A -> B -> C
- Iteration 2: A -> B -> D -> E -> C -> F -> G


Iterative Deepening Search

- Iterative deepening: Try all possible depth bounds in turn.
- Combines depth-first and breadth-first search.
- Does a series of depth-limited depth-first searches to depth 1, 2, 3, etc.
- Early states will be expanded multiple times, but that might not matter too much because most of the nodes are near the leaves.


Iterative Deepening Search



Iterative Deepening Search



Iterative Deepening Search



Uninformed Search

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)

Uniform Cost Search (UCS)

Informed Searches

Greedy Best-First Search

A* Search

Uniform-Cost Search, $f(n) = g(n)$

Sometimes transitions from one node to another have a cost

Cost of a **path** is the sum of the costs of its arcs:

$$\underline{cost}(\langle n_0, \dots, n_k \rangle) = \sum_{i=1}^k cost(\langle n_{i-1}, n_i \rangle)$$

An optimal solution has minimum cost

Delivery robot example:

- cost of arc may be resources (e.g., time, energy) required to execute action represented by the arc
- aim is to reach goal using least resources

Uniform-Cost Search, $f(n) = g(n)$

Expand root first, then expand **least-cost** unexpanded node

- implemented by treating the frontier as a priority queue ordered by the cost function

$$cost(\langle n_0, \dots, n_k \rangle) = \sum_{i=1}^k cost(\langle n_{i-1}, n_i \rangle)$$

Uniform-Cost Search, $f(n) = g(n)$

- Reduces to breadth-first search when all actions have same cost
- Finds the cheapest goal provided path cost is monotonically increasing along each path (i.e. no negative-cost steps)

$$f(n) = g(n)$$

$f(n)$: the estimated total cost through node n

$g(n)$: the cost of the path from the current node to the next node n.

Uninformed Search

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)

Uniform Cost Search (UCS)

Greedy Best-First Search


A* Search

Informed Searches

Uninformed vs Informed Search

- Uninformed - keeps searching until it stumbles on goal
 - No domain knowledge
- Informed - searches (**a.k.a heuristics**) in direction of best guess to goal
 - Uses domain knowledge

Greedy Best-First Search, $f(n) = h(n)$



Greedy Best-First Search, $f(n) = h(n)$

Always select node closest to goal according to heuristic function. $h(n)$ is the estimated cost for a current node to goal.

- $h(n) = 0$ if n is a goal state

Frontier is a **priority queue** ordered by h .

“Greedy” algorithm takes “best” node first.

$$f(n) = h(n)$$

$f(n)$: the estimated total cost through node n

$h(n)$: the **heuristic** estimate of the cost from n to the goal.

Uninformed Search

Search Algorithms Covered at this Lecture

Breath First Search (BFS)

Depth First Search (DFS)

Iterative Deepening Search (IDDFS)


Uniform Cost Search (UCS)

Greedy Best-First Search

A* Search

Informed Searches

A* Search, $f(n) = g(n) + h(n)$



Uniform-Cost Search



Greedy Search



A* Search

A* Search, $f(n) = g(n) + h(n)$

Use both cost of path generated and estimate to goal to order nodes on the frontier

$$f(n) = g(n) + h(n)$$

$f(n)$: the estimated total cost through node n

$g(n)$: the cost of the path from the start node to the current node n.

$h(n)$: the heuristic estimate of the cost from n to the goal.

Order priority queue using function $f(n) = g(n) + h(n)$

Summary of Informed Search

- Informed search makes use of problem-specific knowledge to guide progress of search
- This can lead to a significant improvement in performance
- Greedy Search tries to minimise cost from current node n to the goal.
- A* combines the advantages of Uniform-Cost Search and Greedy Search