

# Reinforcement Learning

COMP3411/9814: Artificial Intelligence

# Lecture Overview

- Introduction
- Elements of Reinforcement Learning
- Exploration vs Exploitation
- The agent-environment interface
- Value functions
- Temporal difference prediction

# Lecture Overview

- **Introduction**
- Elements of Reinforcement Learning
- Exploration vs Exploitation
- The agent-environment interface
- Value functions
- Temporal difference prediction

# Initial ideas

Instrumental or operational conditioning.  
Stimulus-behavior learning.

Thorndike, 1911




# Reinforcement learning


- Idea of learning by interaction with the environment.
- With no explicit instructor but with a direct sensorimotor connection.
- Awareness of how our environment answers to what we do.



# Reinforcement Learning




# Pole balancing



- Pole balancing can be learned the same way
- Reward might be only received at the end
  - after falling or hitting the end of the track

# Pole balancing



# And you think pole balancing is trivial?




Flip Maneuver

# Lecture Overview

- Introduction
- **Elements of Reinforcement Learning**
- Exploration vs Exploitation
- The agent-environment interface
- Value functions
- Temporal difference prediction

# Reinforcement Learning

- RL is to learn what to do, mapping from situations to actions.
- An agent should be able to sense the environment states and perform actions to affect such states.
- Actions might affect not only immediate reward.
- An important challenge is the exploration/ exploitation trade-off problem.



# Elements of reinforcement learning

There are four essential elements:

- **Policy**
  - Informs how to act in a particular situation.
  - Set of stimulus-response rules or associations.
  - Can be stochastic.

# Elements of reinforcement learning

There are four essential elements:

- Reward function
  - Defines the aim of an RL problem.
  - Maps each perceived state (or state-action pair) into a number, the reward.
  - The goal is to maximize the long-term reward.
  - In biological systems may correspond to pain and pleasure feelings.
  - Can be stochastic.

# Elements of reinforcement learning

There are four essential elements:

- **Value function**
  - Shows what it's good in the long run (the reward in an immediate sense).
  - In biological systems corresponds to more refined judgments of foresight about the future from one state.
  - Actions are decided based on the value.
  - It's much harder to determine values than rewards.

# Elements of reinforcement learning

There are four essential elements:

- Optionally, a model of the environment
  - Imitates the environment behaviour.
  - Can predict states and reward obtained.
  - The use of models of the environment is still relatively new.

# Lecture Overview

- Introduction
- Elements of Reinforcement Learning
- **Exploration vs Exploitation**
- The agent-environment interface
- Value functions
- Temporal difference prediction

# Exploration / Exploitation Trade-off

- Most of the time, the agent chooses what it thinks the “best” action is.
- But to learn, it must occasionally choose something different from the preferred action.

# Exploration / Exploitation Trade-off

Should I stay or should I go now?  
Should I stay or should I go now?  
If I go, there will be trouble  
And if I stay it will be double

-- The Clash



# Exploration / Exploitation Trade-off

- The greedy action exploits the current knowledge.
- The non-greedy action explores.
- Exploitation maximises the immediate reward and exploration in the long run.
- There is a conflict between exploration and exploitation.



# Action-value estimation methods

- We denote the real action value as  $q_*(a)$ .
- We denote the estimated value at time-step t as  $Q_t(a)$ .
- **Simple Estimation:** to average received rewards when action  $a$  has been selected  $K_a$  times.

$$Q_t(a) = \frac{R_1 + R_2 + \cdots + R_{K_a}}{K_a}.$$

# Action-value estimation methods

- If  $K_a = 0$ ,  $Q_t(a)$  is defined with an arbitrary value, e.g.,  $Q_t(a) = 0$  (not necessarily the best).
- As  $K_a \rightarrow \infty$ ,  $Q_t(a)$  converges to  $q^*(a)$ .

$$Q_t(a) = \frac{R_1 + R_2 + \cdots + R_{K_a}}{K_a}.$$

# Action-value estimation methods

## Greedy method



- The simplest way to choose an action: the action with the highest estimated value.
- $A_t^*$  where  $Q_t(A_t^*) = \max_a Q_t(a)$ .

## $\epsilon$ -greedy method

- A simple alternative: to choose the best action most of the time, and sometimes (with a small probability  $\epsilon$ ) a random one.
- $Q_t(a)$  converges to  $q_*(a)$  with probability  $1 - \epsilon$ .

# Action-value estimation methods

- 2000 agents averaged.
- It's possible to reduce  $\epsilon$  over time



# Action-value estimation methods

## Softmax method

- $\epsilon$ -greedy effectively trades off exploration and exploitation, but the selection is equitable (or fair) among actions.
- Sometimes, the worst action is very bad.
- High temperatures give almost equal probability for all actions.
- Low temperatures make a bigger difference in the probability.


$$\frac{e^{Q_t(a)/\tau}}{\sum_{i=1}^n e^{Q_t(i)/\tau}}$$

# Lecture Overview

- Introduction
- Elements of Reinforcement Learning
- Exploration vs Exploitation
- **The agent-environment interface**
- Value functions
- Temporal difference prediction


# The agent-environment interface

- Any method able to solve the problem is considered an RL method.
- **Agent:** comprises the learner and the one making the decisions. (although they can be separated!).
- **Environment:** everything external to the agent that it interacts with.




# The agent-environment interface

- Reward: numeric value the agent tries to maximise.  $R_{t+1} \in \mathbb{R}$ .
- $S_t \in S$ .  $S$  set of possible states.
- $A_t \in A(S_t)$ .  $A(S_t)$  actions available at  $S_t$ .
- The agent implements a map from the states toward the action selection probability.
- This is called agent policy  $\pi_t$  where  $\pi_t(a|s)$  is the probability of  $A_t = a$  and  $S_t = s$ .
- RL methods detail how an agent updates its policy as a result of its experience.



# The agent-environment interface

- **Example:** recycling robot.
- The agent decides if (i) actively search for a can, (ii) remains stationary and waits for a can, or (iii) gets back to the home base to recharge (three possible actions).
- The state is determined by the battery state.
- Reward: Mostly zero, positive when collects a can and negative (much higher) when the battery runs empty.




# Goals and rewards

- The agent's goal is to maximise the amount of total reward, not the immediate reward.
- A robot learning to walk receives a reward proportional to the forward movement.
- A robot learning to escape from a maze receives a reward equal to zero until escapes and then receives +1.
- Another strategy is giving a reward of -1 after each movement till escaping.
- An agent learning to play chess receives +1 for winning and -1 for losing.


# Goals and rewards

- The chess player should be rewarded only for winning and not for taking opponent's pieces.
- Otherwise, the agent will learn to maximise subgoals.
- In summary, the reward signal is the way to communicate to the agent **what** you want it to achieve, **not how** you want it achieved.

# Episodic and non-episodic tasks



The pole-balancing task.



The mountain car task.

# Returns

- If the reward sequence received is  $R_{t+1}, R_{t+2}, R_{t+3}, \dots$ . We want to maximise the expected return  $G_t$ .

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T$$

- In tasks with final state and that can be divided into subsequences (episodes)
- Each episode finishes in the final state and the task starts over from an initial state.
- These tasks are known as episodic tasks.

# Returns


- Tasks intended to be performed continuously without limit are referred to as continuous tasks (or non-episodic).
  - The return could be infinite, given that  $T = \infty$ .
  - In this case, the agent maximises the discounted rewards, choosing actions to maximise the discounted return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate  $0 \leq \gamma < 1$ . Determine the present value of future rewards. If  $\gamma = 0$ , the agent is myopic. If  $\gamma \rightarrow 1$  the agent is foresighted.


# Unified Notation

- A final absorbing state with reward equal to zero.



- It is possible that  $T = \infty$  or  $\gamma = 1$ , but not both.

$$G_t = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1}$$



3 minutes

# Returns

| $\gamma$ | Reward sequence        | Return |
|----------|------------------------|--------|
| 0.5      | 1 0 0 0 .....          |        |
| 0.5      | 0 0 2 0 0 0 .....      |        |
| 0.9      | 0 0 2 0 0 0 .....      |        |
| 0.5      | -1 2 6 3 2 0 0 0 ..... |        |

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$




3 minutes

# Returns

| $\gamma$ | Reward sequence        | Return |
|----------|------------------------|--------|
| 0.5      | 1 0 0 0 .....          | 1      |
| 0.5      | 0 0 2 0 0 0 .....      |        |
| 0.9      | 0 0 2 0 0 0 .....      |        |
| 0.5      | -1 2 6 3 2 0 0 0 ..... |        |

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$




3 minutes

# Returns

| $\gamma$ | Reward sequence        | Return |
|----------|------------------------|--------|
| 0.5      | 1 0 0 0 .....          | 1      |
| 0.5      | 0 0 2 0 0 0 .....      | 0.5    |
| 0.9      | 0 0 2 0 0 0 .....      |        |
| 0.5      | -1 2 6 3 2 0 0 0 ..... |        |

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$




4 minutes

# Returns

| $\gamma$ | Reward sequence        | Return |
|----------|------------------------|--------|
| 0.5      | 1 0 0 0 .....          | 1      |
| 0.5      | 0 0 2 0 0 0 .....      | 0.5    |
| 0.9      | 0 0 2 0 0 0 .....      | 1.62   |
| 0.5      | -1 2 6 3 2 0 0 0 ..... |        |

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



3 minutes

# Returns

| $\gamma$ | Reward sequence        | Return |
|----------|------------------------|--------|
| 0.5      | 1 0 0 0 .....          | 1      |
| 0.5      | 0 0 2 0 0 0 .....      | 0.5    |
| 0.9      | 0 0 2 0 0 0 .....      | 1.62   |
| 0.5      | -1 2 6 3 2 0 0 0 ..... | 2      |

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

# The Markov property

- In RL, state means any information available for the agent (either processed or not).
- The state should not inform everything about the environment to the agent. For instance, an agent playing blackjack should not know the next card in the deck.
- We do not blame the agent for not knowing something important, but we do for knowing something and then forgetting it.
- Ideally, a state should contain compact information about the past, retaining relevant information. This is called the Markov property. For instance, the chess board.

# The Markov property

- Sometimes, the property cannot be fully satisfied.
- In pole-balancing, the state satisfies the property if the exact position and velocity of the cart are specified, and the pole angle and its change rate.
- However, there may exist distortions, such as delays and other effects as the temperature of the wheels.
- Some studies have even used a simple region division: left, right, and centre.

# Markov Decision Processes

- An RL task with the Markov property is called Markov decision process (MDP).
- Markov decision process (MDP):  $\langle S, A, \delta, r \rangle$ .
  - $S$  is a finite set of states,
  - $A$  is a set of actions,
  - $\delta$  is the transition function  $\delta: S \times A \rightarrow S$ , and,
  - $r$  is the reward function  $r: S \times A \rightarrow \mathbb{R}$ .

# Recycling robot MDP

- At each moment, the robot decides if (i) actively search for a can, (ii) waits for someone to bring a can, or (iii) gets back to the home base to recharge.
- The best strategy is to actively search for cans.
- In case the battery runs out, the robot needs to be rescued leading to a negative reward.
- The agent solely decides as a function of the energy level of the battery. Two levels: high, low.
- $S = \{\text{high}, \text{low}\}$ .
- Possible decisions (agent's actions): wait, search, recharge.
- $A(\text{high}) = \{\text{search}, \text{wait}\}$ .
- $A(\text{low}) = \{\text{search}, \text{wait}, \text{recharge}\}$ .

# Recycling robot MDP


- Transition probabilities and expected reward:

| $s$  | $s'$ | $a$      | $p(s' s, a)$ | $r(s, a, s')$       |                                       |
|------|------|----------|--------------|---------------------|---------------------------------------|
| high | high | search   | $\alpha$     | $r_{\text{search}}$ |                                       |
| high | low  | search   | $1 - \alpha$ | $r_{\text{search}}$ |                                       |
| low  | high | search   | $1 - \beta$  | -3                  |                                       |
| low  | low  | search   | $\beta$      | $r_{\text{search}}$ | $r_{\text{search}} > r_{\text{wait}}$ |
| high | high | wait     | 1            | $r_{\text{wait}}$   |                                       |
| high | low  | wait     | 0            | $r_{\text{wait}}$   |                                       |
| low  | high | wait     | 0            | $r_{\text{wait}}$   |                                       |
| low  | low  | wait     | 1            | $r_{\text{wait}}$   |                                       |
| low  | high | recharge | 1            | 0                   |                                       |
| low  | low  | recharge | 0            | 0.                  |                                       |

- Assumption: cans cannot be collected when going back to the home base or when the battery is depleted.

# Recycling robot MDP

- Transition graph:



- Transition probabilities from one action always sum to 1.

# Lecture Overview

- Introduction
- Elements of Reinforcement Learning
- Exploration vs Exploitation
- The agent-environment interface
- **Value functions**
- Temporal difference prediction

# Value function

- Value function estimations.
  - State-value function, or
  - Action-value function (for state-action pairs)
- The function estimates how good it is for the agent to be in a given state, in terms of future reward (or expected return).
- The value of a state  $s$  under a policy  $\pi$ , denoted  $v_\pi(s)$  or  $V^\pi(S)$ , is the expected return when starting in  $s$  and following  $\pi$  thereafter:

$$v_\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s] = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right]$$

# Value function

- The value of a terminal state, if any, is zero.
- The value of taking action  $a$  in state  $s$  under policy  $\pi$ , is denoted  $q_\pi(s,a)$  or  $Q^\pi(S,A)$ :

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a] = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right]$$

- Value functions  $v_\pi(s)$  and  $q_\pi(s,a)$  can be estimated from experience.
- If there are many states, it's impractical to keep values for each state.
- In this case, parameterized function approximators are used to keep  $v_\pi(s)$  and  $q_\pi(s,a)$ .

# Value function

- Try to maximise expected future reward:

$$\begin{aligned} V^\pi(s_t) &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\ &= \sum_{i=0}^{\infty} \gamma^i r_{t+i} \end{aligned}$$

- $V^\pi(s_t)$  is the value of state  $s_t$  under policy  $\pi$
- $\gamma$  is a discount factor [0..1]

# Value function


- $V^\pi(s)$  is the expected value of following policy  $\pi$  in state  $s$
- $V^*(s)$  be the maximum discounted reward obtainable from  $s$ .
  - i.e. the value of following the optimal policy
- We make the simplification that actions are deterministic, but we don't know which action to take.
  - Other RL algorithms relax this assumption

# Value function

- The red arrows show  $\pi^*$ , the optimal policy, with  $\gamma = 0.9$
- $V^*(s)$  values shown in red

$$V^\pi(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

$$= \sum_{i=0}^{\infty} \gamma^i r_{t+i}$$




# $Q$ -values

- How to choose an action in a state?


$$Q(s, a) = r(s, a) + \gamma V^*(s')$$

- The  $Q$ -value for an action,  $a$ , in a state,  $s$ , is the immediate reward for the action plus the discounted value of following the optimal policy after that action
- $V^*$  is value obtained by following the optimal policy
- $s' = \delta(s, a)$  is the succeeding state, assuming the optimal policy


# $Q$ -values



$r(s, a)$  (immediate reward) values




$V^*(s)$  values





$Q(s, a)$  values

$$\gamma = 0.9$$

# Grid world example





# Another grid world example



|      |      |      |      |      |
|------|------|------|------|------|
| 3.3  | 8.8  | 4.4  | 5.3  | 1.5  |
| 1.5  | 3.0  | 2.3  | 1.9  | 0.5  |
| 0.1  | 0.7  | 0.7  | 0.4  | -0.4 |
| -1.0 | -0.4 | -0.4 | -0.6 | -1.2 |
| -1.9 | -1.3 | -1.2 | -1.4 | -2.0 |

- Cells correspond to the states.
- 4 possible actions.
- Actions leading the agent out of the environment do not change the position but give reward = -1.
- All other actions give reward = 0, except movements from A and B.



# Another grid world example



|      |      |      |      |      |
|------|------|------|------|------|
| 3.3  | 8.8  | 4.4  | 5.3  | 1.5  |
| 1.5  | 3.0  | 2.3  | 1.9  | 0.5  |
| 0.1  | 0.7  | 0.7  | 0.4  | -0.4 |
| -1.0 | -0.4 | -0.4 | -0.6 | -1.2 |
| -1.9 | -1.3 | -1.2 | -1.4 | -2.0 |


- All actions equal probability.
- Discount factor  $\gamma = 0.9$ .
- Negative values near the lower edge.
- The best state is A, but expected return is lower than 10, the immediate reward.
- B is valued more than 5, the immediate rewards.

# Another grid world example




|      |      |      |      |      |
|------|------|------|------|------|
| 3.3  | 8.8  | 4.4  | 5.3  | 1.5  |
| 1.5  | 3.0  | 2.3  | 1.9  | 0.5  |
| 0.1  | 0.7  | 0.7  | 0.4  | -0.4 |
| -1.0 | -0.4 | -0.4 | -0.6 | -1.2 |
| -1.9 | -1.3 | -1.2 | -1.4 | -2.0 |

- The best state is A, but expected return is lower than 10, the immediate reward.
- B is valued more than 5, the immediate rewards.
- Why?



2 minutes


# Another grid world example



a) gridworld

|      |      |      |      |      |
|------|------|------|------|------|
| 22.0 | 24.4 | 22.0 | 19.4 | 17.5 |
| 19.8 | 22.0 | 19.8 | 17.8 | 16.0 |
| 17.8 | 19.8 | 17.8 | 16.0 | 14.4 |
| 16.0 | 17.8 | 16.0 | 14.4 | 13.0 |
| 14.4 | 16.0 | 14.4 | 13.0 | 11.7 |

b)  $v_*$



c)  $\pi_*$

- Optimal value function and optimal policy for the grid world.

# Lecture Overview

- Introduction
- Elements of Reinforcement Learning
- Exploration vs Exploitation
- The agent-environment interface
- Value functions
- Temporal difference prediction

# Temporal-difference (TD) prediction

- TD is one central and novel idea in RL.
- Monte Carlo-like methods must wait until the end of the episode to update  $V(S_t)$  – (only at that point  $G_t$  is known):


$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

- The simplest TD method is called TD(0):

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

# Temporal-difference (TD) prediction

- Approximations can be on-policy or off-policy.
- TD control learns an action-value function instead of a state-value function.
- We estimate  $q_{\pi}(s,a)$  for the current policy  $\pi$ .
- Therefore, we consider transitions from state-action pair to state-action pair.



# Sarsa: On-Policy TD Control

- Updates after each transition from a non-terminal  $S_t$ .
- If  $S_{t+1}$  is terminal,  $Q(S_{t+1}, A_{t+1})$  is defined as zero.

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- Each element of the 5-tuple  $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$  is used, this gives the name to the algorithm.
- On-policy methods continuously estimate  $q_\pi$  for policy  $\pi$ , and at the same time change  $\pi$  greedily towards  $q_\pi$ .

# Sarsa: On-Policy TD Control

- On-policy TD algorithm:

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Repeat (for each step of episode):

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$ ;

    until  $S$  is terminal

# Q-Learning: Off-Policy TD Control

- A simple but important breakthrough is an off-policy TD algorithm.
- The simplest way is *one-step Q-learning*:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

- The learned action-value function  $Q$  directly approximates  $q^*$ , the optimal action-value function, regardless the followed policy  $\pi$ .
- The policy still has an effect in which state-action pairs are visited and updated.

# Q-Learning: Off-Policy TD Control

- Off-policy TD algorithm:

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

    Initialize  $S$

    Repeat (for each step of episode):

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$


$$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$ ;

    until  $S$  is terminal



# The Cliff Walking

- Reward of -1 for all transitions, except in the cliff.
- The cliff gives a negative reward of -100 and sends the agent back to the start position.




# The Cliff Walking

- $\epsilon$ -greedy, with  $\epsilon=0.1$ .
- Q-learning learns the optimal path. Sarsa learns the longest, safest path.
- However, overall Q-learning behaviour is worse.
- If  $\epsilon$  is gradually reduced, both methods converge asymptotically to the optimal policy.




# Deep Q-Network

- Proposed by Mnih et al. in 2015.
- Human-level control through deep reinforcement learning.
- Tested in 49 Atari games.
- ANN to estimate values.
- An additional target network is used for stability.
- Use a memory buffer for experience replay.




# Actor-critic methods

- Policy approximation.
- Learning is always on-policy.
- The actor structures the policy.
- The critic must learn and critique the followed policy.
- Minimal computation to select actions, even in continuous-valued actions or for stochastic policies.
- The separate actor in actor-critic is more appealing as psychological and biological models.



# Taxonomy of RL Algorithms


- Non-exhaustive modern taxonomy
- From OpenAI Spinning up library



- [https://spinningup.openai.com/en/latest/spinningup/rl\\_intro2.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro2.html)


# Examples

- Stumpy - A simple learning robot.
- Stumpy receives a *reward* after each action. Did it move forward or not?
- After each move, updates its policy.




# Examples

- Stumpy - A simple learning robot.
- Continues trying to maximise its reward.
- Stumpy after 30 minutes.




# Examples

- Another example



# Reference

- For a more comprehensive introduction, you should definitely have a look at:
  - Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.
  - <http://www.incompleteideas.net/book/the-book-2nd.html>



# Feedback

- In case you want to provide anonymous feedback on these lectures, please visit:
- <https://forms.gle/KBkN744QuffuAZLF8>

Muchas gracias!



AI Lecture Feedback

This is a short form to provide early feedback for lectures

franciscocruzhh@gmail.com [Switch account](#) 

✉ Not shared

\* Indicates required question

In case you want a reply, provide your zID. Otherwise your answer is anonymous.

Your answer

how did you participate? \*

In the classroom

Watch the class from automatic recording

If you have any comments, feedback, or question about the lectures, this is the place.

Your answer

[Submit](#) [Clear form](#)